

No cascading-rollbacks!

Overview over this video

Last video: Recoverable schedules

- I.e. if we follow those, we can avoid breaking durability (because of cascading rollbacks)

This video: Cascadeless schedules

- I.e. if we follow those, we can avoid cascading rollbacks!

Recoverable schedules still cascades

A recoverable schedule:

$S_4: w_1(X); w_1(Y); w_2(X); r_2(Y); w_2(Y); c_1; c_2$



Suppose T_1 needs to be rolled back here

T_1 rolls back $\rightarrow T_2$ has to be rolled back


Cascadeless Schedules

读已提交

A schedule is **cascadeless** if each transaction in it reads only values that were written by transactions that have already committed.

读已提交
所以只读
①

No reading of "dirty data".
No cascading rollbacks.



Cascadeless Schedules

A schedule is **cascadeless** if each transaction in it reads only values that were written by transactions that have already committed.

No reading of “dirty data”.
No cascading rollbacks.



As for recoverable schedules:
Log records have to reach disk in the right order.

Example

The schedules S_1 - S_4 in previous video are **not cascadeless**:

S_1 : $w_2(X); w_1(Y); w_1(X); r_2(Y); w_2(Y); c_1; c_2$

S_2 : $w_1(X); w_1(Y); w_2(X); r_2(Y); w_2(Y); c_2; c_1$

S_3 : $w_2(X); w_1(Y); w_1(X); r_2(Y); w_2(Y); c_2; c_1$

S_4 : $w_1(X); w_1(Y); w_2(X); r_2(Y); w_2(Y); c_1; c_2$

reads uncommitted data from T_1

This variant of S_1 is **cascadeless**:

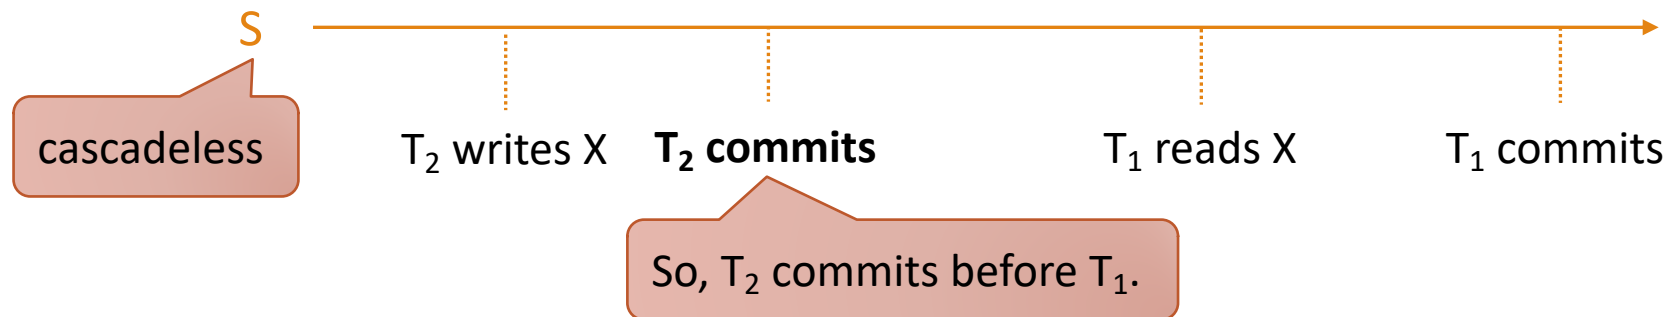
S_5 : $w_2(X); w_1(Y); w_1(X); c_1; r_2(Y); w_2(Y); c_2$

Note: S_5 is **not serialisable**.

reads committed data from T_1

Cascadeless Schedules: Properties

Cascadeless schedules are **recoverable**:

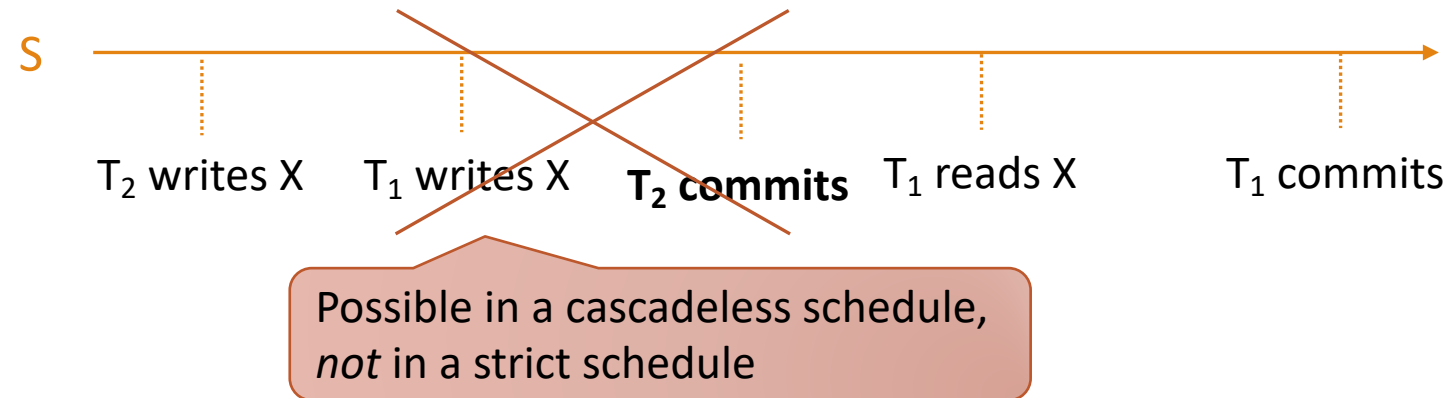


Cascadeless schedules are in general **not serialisable**.
(recall example on previous slide)

Can We Have Both?
No Cascading Aborts & Serialisability?

Strict Schedules

A schedule is **strict** if each transaction in it reads and writes only values that were written by transactions that have already committed.



Strict Schedules

A schedule is **strict** if each transaction in it reads and writes only values that were written by transactions that have already committed.

把已提交的数据



Of course, log records have to reach disk in order.

Strict Two-Phase Locking (Strict 2PL)

Most popular variant of two-phase locking (2PL)

Enforces both:

- Conflict-serialisability
- Strict schedules

Strict locking condition
(in addition to 2PL condition):

- with simple locking: **any lock**
- with shared/exclusive locks: **just exclusive locks**

A transaction T **must not release any lock (that allows T to write data)** until:

- T has committed or aborted, and
- the commit/abort log record has been written to disk.

Example 1

2PL transaction

Not strict 2PL

Transaction T

lock(X)
read_item(X)
X := X + 100
write_item(X)
lock(Y)
unlock(X)
read_item(Y)
Y := Y + 100
write_item(Y)
unlock(Y)
commit

For undo logging, we assume that **commit...**

1. Writes all log records to disk
2. Writes all modified database items to disk
3. Writes the commit record to disk

Example 2

2PL transaction

Strict 2PL transaction

Transaction T

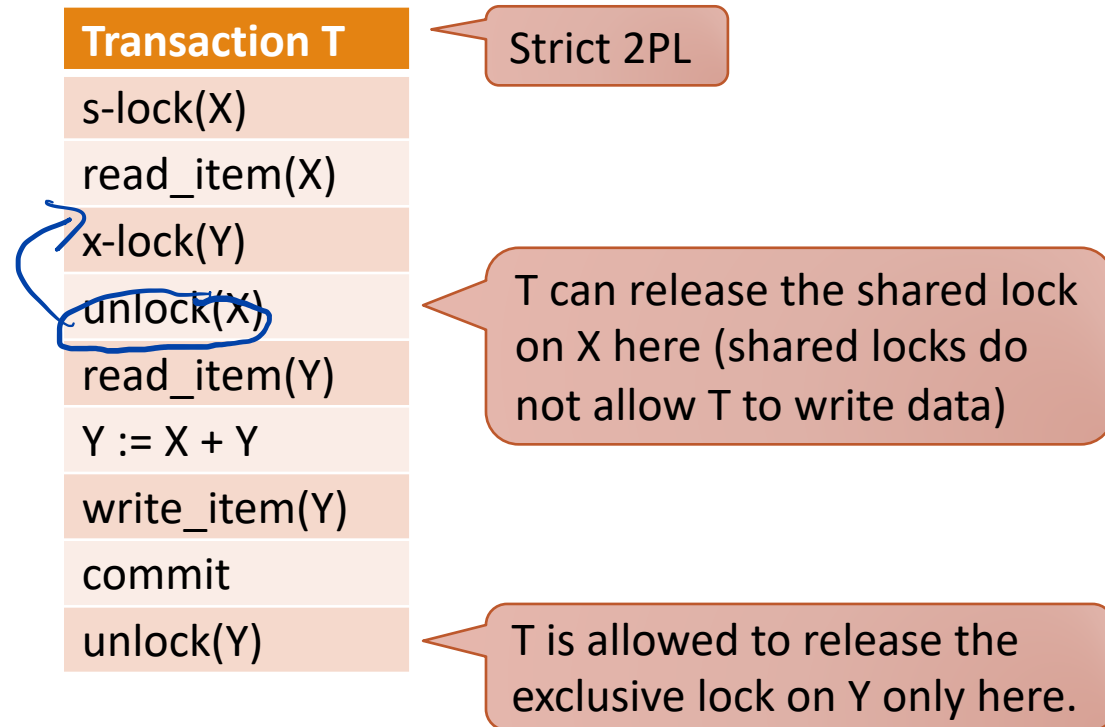
lock(X)
read_item(X)
X := X + 100
write_item(X)
lock(Y)
unlock(X)
read_item(Y)
Y := Y + 100
write_item(Y)
unlock(Y)
commit

New transaction T'

lock(X)
read_item(X)
X := X + 100
write_item(X)
lock(Y)
read_item(Y)
Y := Y + 100
write_item(Y)
commit
unlock(X)
unlock(Y)

Locks released only after fully committed,
and all log records written to disk

Example With Shared/Exclusive Locks

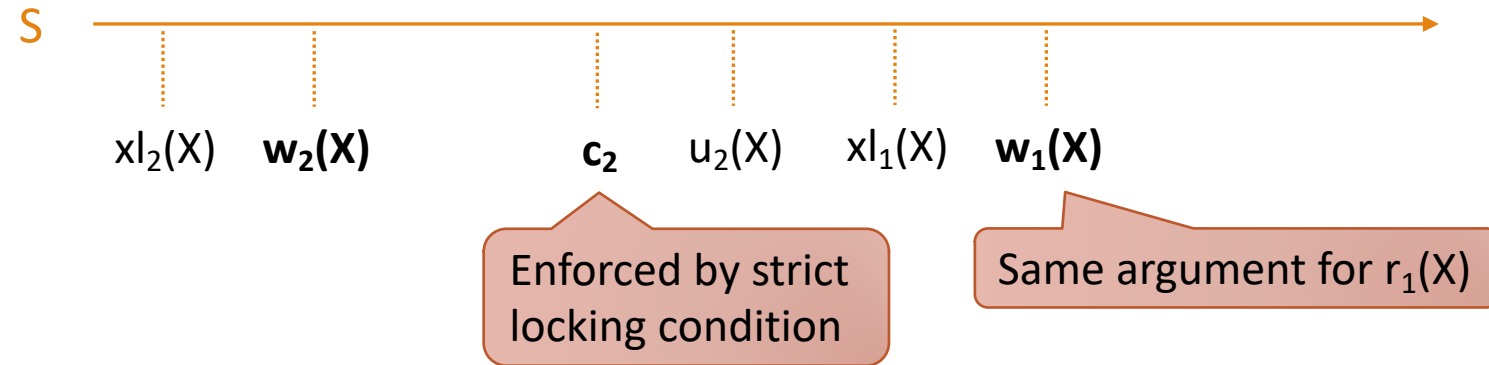


Strict Two-Phase Locking *Enforces* Conflict-Serialisable & Strict Schedules

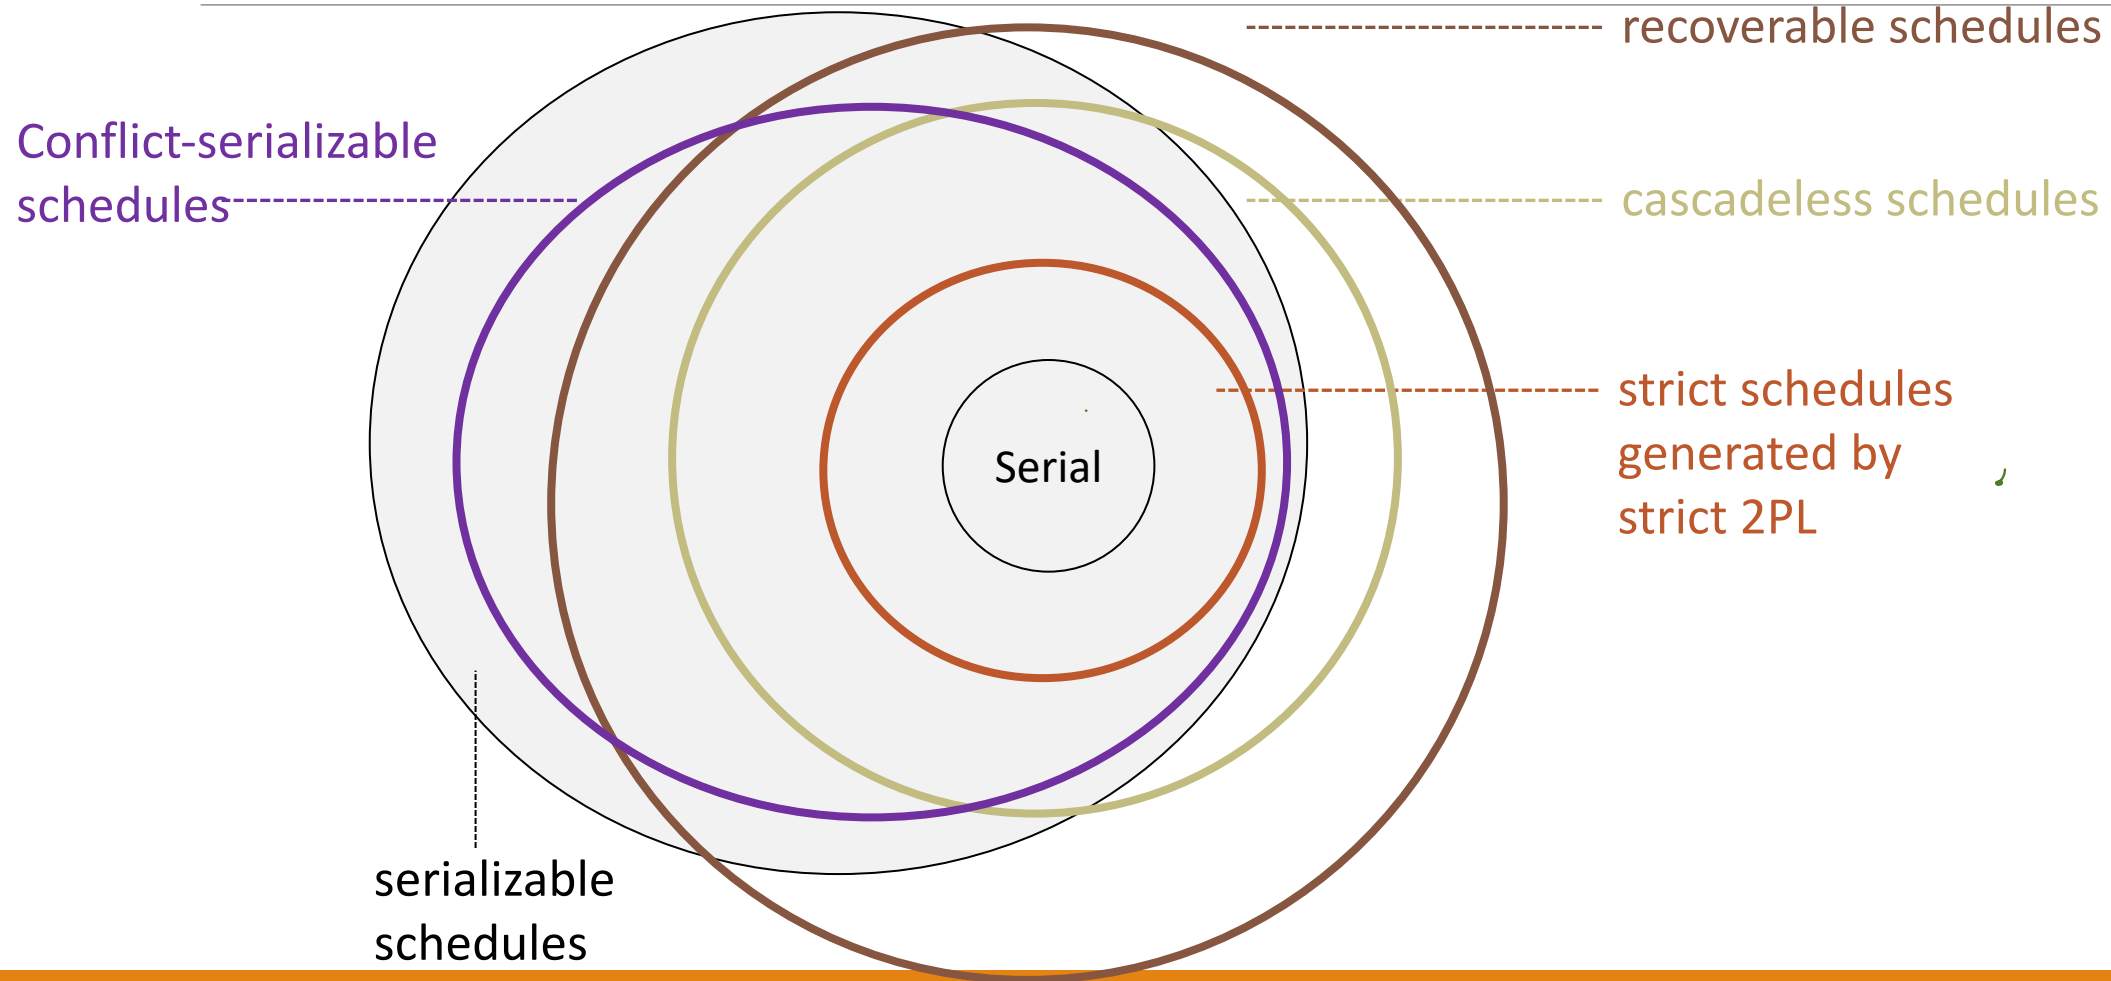
If S is a schedule consisting of strict 2PL transactions:

- S is conflict-serialisable.
- S is strict.

Strictness:



How the Types of Schedules are Related



Examples for serializable schedules

都是对X修改，互不影响。

1. example: $w_2(X); w_1(X); w_1(Y); w_2(Y); r_3(Y); w_3(X); c_3; c_2; c_1$

- This is serializable, but not conflict-serializable or recoverable

2. example: $w_1(X); w_1(Y); w_2(X); w_2(Y); r_3(Y); w_3(X); c_3; c_2; c_1$

- This is conflict-serializable, but not recoverable

3. example: $w_2(X); w_1(X); w_1(Y); w_2(Y); r_3(Y); w_3(X); c_2; c_3; c_1$

- This is recoverable and serializable, but not conflict-serializable nor cascadeless

4. example: $w_1(X); w_1(Y); w_2(X); w_2(Y); r_3(Y); w_3(X); c_2; c_3; c_1$

- This is recoverable and conflict-serializable, but not cascadeless

5. example: $w_2(X); w_1(X); w_1(Y); w_2(Y); c_2; r_3(Y); w_3(X); c_3; c_1$

- This is cascadeless and serializable, but not conflict-serializable

6. example: $w_1(X); w_1(Y); w_2(X); w_2(Y); c_2; r_3(Y); w_3(X); c_3; c_1$

Conflict-serializable
and cascadeless but
not strict

fx-14p27 write 1454

1/18 2025/5 16 commit 14

Still... Risk of Deadlocks

T ₁
lock(X)
read_item(X)
X := X + 100
write_item(X)
lock(Y)
read_item(Y)
Y := Y + 100
write_item(Y)
commit
unlock(X)
unlock(Y)

T ₂
lock(Y)
read_item(Y)
Y := Y + 100
write_item(Y)
lock(X)
read_item(X)
X := X + 100
write_item(X)
commit
unlock(X)
unlock(Y)

T₂'s request for lock on X denied

$l_1(X); r_1(X); w_1(X); l_2(Y); r_2(Y); w_2(Y); \underline{\quad ? \quad}$

T₁'s request for lock on Y denied

Strict 2PL and Deadlocks

Strict 2PL yields **conflict-serialisable, strict schedules**

Problem: **deadlocks**

T ₁	T ₂
lock(X)	lock(Y)
read_item(X)	read_item(Y)
X := X + 100	Y := Y + 100
write_item(X)	write_item(Y)
lock(Y)	lock(X)
...	...

$l_1(X); r_1(X); w_1(X);$
 $l_2(Y); r_2(Y); w_2(Y); \underline{\quad ? \quad}$

Roll back (and restart)
one of the transactions

Two approaches for deadlock prevention:

- **Detect deadlocks & fix them**
- **Enforce deadlock-free schedules**

Not based on (strict) 2PL

Summary

Cascadeless schedules are a more restrictive form of schedules

- You can only read things that are already committed

Instead of implementing that, we use Strict

- You can only read or overwrite things that are already committed

We then finally use Strict 2PL

- Like 2PL, but where you can only unlock (locks that could write) after commit