

Schedules, a low-level view and notation

Overview of this video

The video will introduce a simplified and lower level view of transactions as well as schedules and notations for all of these

Translating SQL into Low-Level Operations

Employees(e_id, first_name, last_name, birth_day, salary)

Two SQL Statements

```
SELECT salary
FROM Employees
WHERE e_id = 1234;
```

```
UPDATE Employees
SET salary = salary*1.1
WHERE e_id = 1234;
```

produces



THREE TRANSACTION OPERATIONS

1. read(e_id=1234, salary);
2. salary=salary*1.1;
3. write(e_id=1234, salary);

Notes:

Abstraction (at a low level)

Read data item 'salary' from tuple with primary key 1234

Two database operations: op1 (read) and op3 (write)

One non-database operation: op2 (the calculation)

Simplifying Low-level Operations

THREE TRANSACTION OPERATIONS

1. `read(e_id=1234, salary);`
2. `salary=salary*1.1;`
3. `write(e_id=1234, salary);`

produces



THREE SIMPLIFIED TRANSACTION OPERATIONS

1. `read(X);`
2. `X=X*1.1;`
3. `write(X);`

X is in this case e_id 1234's salary, but it does not really matter and we omit it

Basic Operations of Transactions

read(X): Reads a database item X into a program variable (also named X, for simplicity)

- Find the address of the disk block (page) that contains item X
- Copy that disk block into a buffer in main memory
 - if that disk block is not already in some main memory buffer
- Copy item X from the buffer to the program variable X

write(X): Writes the value of program variable X into the database item named X

- Find the address of the disk block (page) that contains item X.
- Copy that disk block into a buffer in main memory
 - if that disk block is not already in some main memory buffer.
 - Copy item X from the program variable X into its correct location in the buffer
 - Store the updated block from the buffer back to disk either immediately or at some later point in time.

Transactions (in general)

A logical unit of processing using access operations

- Begin
- End
- read(retrieval – **SELECT** etc.)
- write(**insert, update, or delete**)
- + other non-database operations

Begin/end are omitted
when the beginning/end
of a transaction are understood

Schedules

Schedules hold many transactions for execution

The operations making up the transactions are then executed by the schedule in some order

- It **must** preserve that the operations in each transaction happens in the right order!

Two types:

- **Serial Schedules**
 - Executes the transactions one after another (i.e. first each operation of the first schedule, then each operation of the second and so on)
- **Concurrent Schedules**
 - Can interleave operations from the transactions (while still preserving that the operations in each transaction happens in the right order) - formally speaking, a serial schedule is therefore also concurrent...

A Serial Schedule

Executes all operations in transaction **T1**,
then all operations in transaction **T2**.

For simplicity we will typically ignore
the non-database operations...

Begin (T1)

read(X);

$X := X + 100;$

write(X);

read(Y);

$Y := Y + 50;$

write(Y);

commit;

End (T1)

Begin (T2)

read(X);

read(Y);

$X := X + Y;$

write(X);

commit;

End (T2)

Shorthand Notation for Schedules

Shorthand notation for this schedule:

S_a : $r_1(X); w_1(X); r_1(Y); w_1(Y); c_1; r_2(X); r_2(Y); w_2(X); c_2$

Symbols:

- S_{id} = schedule (id is the schedule ID)
- $r_i(X)$ = read(X) in transaction i
- $w_i(X)$ = write(X) in transaction i
- c_i = commit in transaction i
- a_i = abort ("rollback") in transaction i

```
Begin (T1)
  read(X);
  X := X + 100;
  write(X);
  read(Y);
  Y := Y + 50;
  write(Y);
  commit;
```

End (T1)

```
Begin (T2)
  read(X);
  read(Y);
  X := X + Y;
  write(X);
  commit;
```

End (T2)

Another Example

Time	S_b	X
t0		100
t1	read(X)	100
t2	$X = X - 10$	90
t3	write(X)	90
t4	Commit	90
t5	read(X)	90
t6	$X = X * 10$	900
t7	write(X)	900
t8	commit	900

What is the shorthand notation for this schedule?

Another Example

Time	S_b	X
t0		100
t1	read(X)	100
t2	$X = X - 10$	90
t3	write(X)	90
t4	Commit	90
t5	read(X)	90
t6	$X = X * 10$	900
t7	write(X)	900
t8	commit	900

What is the shorthand notation for this schedule?

S_b : $r_1(x)$; $w_1(x)$; c_1 ; $r_2(x)$; $w_2(x)$; c_2

Order matters:

Time	S_b	X
t0		100
t1	read(X)	100
t2	$X = X - 10$	90
t3	write(X)	90
t4	Commit	90
t5	read(X)	90
t6	$X = X * 10$	900
t7	write(X)	900
t8	commit	900

vs.

Time	S_c	X
t0		100
t1	read(X)	100
t2	$X = X * 10$	1000
t3	write(X)	1000
t4	Commit	1000
t5	read(X)	1000
t6	$X = X - 10$	990
t7	write(X)	990
t8	commit	990

Concurrent Schedule

Shorthand notation for schedule:

$S_a: r_1(X); w_1(X); r_1(Y); w_1(y); c_1; r_2(X); r_2(Y); w_2(X); c_2$

$S_d: r_2(X); r_2(Y); w_2(X); c_2; r_1(X); w_1(X); r_1(Y); w_1(y); c_1$

Note that these are **serial schedules** as well as a **concurrent schedules** (in that **all serial schedules are concurrent schedule!**)

Examples of other concurrent schedules (that are not serial):

$S_e: r_1(X); w_1(X); r_2(X); r_1(Y); w_1(y); c_1; r_2(Y); w_2(X); c_2$

$S_f: r_1(X); r_2(X); w_1(X); r_2(Y); r_1(Y); w_2(X); w_1(y); c_2; c_1$

Examples of something that is **not** a schedule:

$S_g: r_1(X); r_2(Y); w_1(X); r_2(X); r_1(Y); w_2(X); w_1(y); c_2; c_1$

Begin (T1)

read(X);
X := X + 100;
write(X);
read(Y);
Y := Y + 50;
write(Y);
commit;

End (T1)

Begin (T2)

read(X);
read(Y);
X := X + Y;
write(X);
commit;

End (T2)

Summary

We want to focus on low-level details (i.e. reads and writes) and their interaction over high-level (i.e. queries) when we talk about transactions (which is a sequence of queries) and schedules (which is a set of transactions that should be executed, so that each operation in each transaction comes in order)

We have a notations for such, using symbols:

- S_{id} = schedule (id is the schedule ID)
- $r_i(X)$ = read(X) in transaction i
- $w_i(X)$ = write(X) in transaction i
- c_i = commit in transaction i
- a_i = abort (“rollback”) in transaction i